

HOMEWORK 4

11-967

Due date: 02/16/2025 11:59 PM EST

In this homework, you will learn how to tokenize data.

Problem 1: BPE Tokenizer

Formally, a tokenizer is a function that takes a string $s \in \mathcal{A}^*$ in alphabet \mathcal{A} to a sequence of tokens $\mathbf{x} \in \mathcal{V}^*$ in vocabulary \mathcal{V} .¹ By definition, tokenizers are complete in their alphabet: *any string over \mathcal{A} can be tokenized*. Tokenizers are essential in NLP pipelines because they convert natural language to symbols that machines can parse, and vice versa.

[Question 1.1] (Writing, 12 points) We may expect a tokenizer to satisfy the properties below:

- **Injective:** for all $s \neq s'$, $\text{tokenize}(s) \neq \text{tokenize}(s')$.
- **Invertible:** for any string s , $s = \text{detokenize}(\text{tokenize}(s))$.
- **Preserves concatenation:** for two strings s, s' , $\text{tokenize}(s + s') = \text{tokenize}(s) + \text{tokenize}(s')$.

However, all three properties are violated by the tokenizers of multiple popular language models.

DELIVERABLES FOR Q1.1

Answer the following questions.

- Provide counterexamples showing how a popular, publicly-available tokenizer fails each of the above properties. Write code in `tests/test_tokenizer.py` that causes the assertions in `test_not_injective()`, `test_not_invertible()` and `test_not_preserving_concat()` to all pass. Then, in your report, write down no more than two sentences about each test, describing the tokenizer you chose and the violation you identified.
- Comment on why it is generally impossible to build non-trivial tokenizers that preserve concatenation.^a

^aA trivial tokenizer is one that uses the alphabet as vocabulary.

[Question 1.2] (Coding, 25 points) Byte-pair encoding (BPE) is a technique of *learning* a useful tokenizer from a text corpus. BPE first initializes the set of tokens as all characters in some alphabet, and iteratively merges the most frequent pair of adjacent tokens (bigram) as a new token. Each iteration adds a new token to the vocabulary while retaining the existing tokens.

DELIVERABLES FOR Q1.2

Make the following modifications to the starter code.

- Implement the BPE algorithm by filling in `ASCIIBPETokenizer.merge()` in `src/tokenizer/bpe.py`. Your implementation should pass all of the `test_*_merge()` test cases in `tests/test_tokenizer.py`.
- Now, implement `encode()` and `decode()` functions in `src/tokenizer/bpe.py`. Your implementation should pass `test_encode()` and `test_decode()`.

Additional automatic tests may be performed after submission.

¹For example, the alphabet \mathcal{A} could be the ASCII or the Unicode character set.

[Question 1.3] (*Writing, 6 points*) You have implemented an ASCII tokenizer that handles the [ASCII characters](#). One limitation of ASCII is that it does not include non-Latin characters, for example those of Hindi or Chinese. In contrast, the Unicode standard supports characters from all writing systems.

The instructors have provided a unicode BPE tokenizer code implementation, `src/tokenizer/bpe.py`, and a tokenizer file, `data/english-tokenizer.json`, that has been trained only on English text. Run `python src/tokenizer/visualize.py` to load in this tokenizer and observe the learned tokens.

DELIVERABLES FOR Q1.3

Answer the following questions with at most two sentences each:

- A. What are the longest token that you see, and what does this tell you about the training data?
- B. How can BPE tokenization compromise the privacy of its training data?

[Question 1.4] (*Writing, 6 points*) Try to tokenize a piece of provided English text and its translation in Thai by running `python src/tokenizer/multilingual.py`.

DELIVERABLES FOR Q1.4

Answer the following questions with at most two sentences each:

- A. Is there a big difference between the number of tokens used to tokenize the same document in English and Thai? Explain why.
- B. How could this phenomenon be problematic for users of low-resource languages?